

# Lecture 9: Algorithm building blocks and CUDA Libraries

## Informatik elective: GPU Computing

Pratik Nayak

Licensed under



# In this session

- Basic data structures, C++ standard library equivalents for the GPU → thrust
- BLAS operations → cuBLAS
- Sparse functionality → cuSPARSE
- Tensors → cuTENSOR
- Solvers → cuSOLVER

# C++ STL on GPUs

- The C++ standard template library (STL) contains many useful data structures, containers, and primitive operations on these containers which can be used to develop more complex algorithms.
- Thrust ( <https://nvidia.github.io/cccl/thrust> ) is the C++ template library for CUDA. It is based on STL.
- It provides similar containers and primitive operations as STL, with only slight changes in interface, adapting for CUDA GPUs

# C++ STL on GPUs

- Similar to `std::vector`, you have `thrust::host_vector` and `thrust::device_vector`
  - Dynamically resizable generic templated containers.
- Operations such as `thrust::fill`, `thrust::copy`, `thrust::sequence` are available.
- Copy data from host to device and back with overloaded “=” operator.
  - `dev_vec = host_vec` // Copies data from host to device with `cudaMemcpy`
- Generic operations with lambdas using `thrust::transform`, in `<thrust/functional>`.
- For example, a simple axpy looks like

# Axpy with thrust

```
struct saxpy_functor{
    const float a;
    saxpy_functor(float _a) : a(_a) {}
    __host__ __device__ float operator()(const float& x, const float& y) const {
        return a * x + y;
    }
};

void saxpy_fast(float A, thrust::device_vector<float>& X, thrust::device_vector<float>& Y){
    // Y <- A * X + Y
    thrust::transform(X.begin(), X.end(), Y.begin(), Y.begin(), saxpy_functor(A));
}
```

# Reduction

- Reduction operations are also available: `thrust::reduce`, `thrust::count`, `thrust::count_if` etc.
- Sum of elements in a vector: `auto sum = thrust::reduce(vec.begin(), vec.end())`
- You can better performance and more idiomatic C++ code by fusing kernels with `thrust::transform_reduce`
  - Better optimization from runtime and more favorable cache behaviour.

# L2-norm using thrust

```
template<typename T>
struct square{
    __host__ __device__ T operator()(const T& x) const {
        return x * x;
    }
};

template<typename T>
auto norm(thrust::device_vector<T>& X){
    // return <- sqrt(sum(X * X))
    return std::sqrt(thrust::transform_reduce(X.begin(), X.end(), square<T>{}, T{0},
        thrust::plus<T>{}));
}
```

# Scans, reorderings and sorting on GPU data

- Index operations generally make use of scan operations: inclusive, exclusive scans:
  - $\{1, 0, 2, 2, 1, 3\} \rightarrow \text{inclusive\_scan} \rightarrow \{1, 1, 3, 5, 6, 9\}$
- Reordering operations to select/copy elements satisfying some conditions:
  - `copy_if`: copy elements that satisfy a condition
  - `partition`: reorder elements based on a condition
  - `remove`, `remove_if`: remove elements
  - `unique`: remove consecutive duplicates
- Sorting operations
  - `sort` and `stable_sort`: equivalent behaviour as in STL.
  - Additionally `sort_by_key` enables sorting of key-value pairs



# Scans, reorderings and sorting on GPU data

- ```
#include <thrust/sort.h>
```
- ```
...  
const int N = 6;  
int    keys[N] = { 1,  4,  2,  8,  5,  7};  
char values[N] = {'a', 'b', 'c', 'd', 'e', 'f'};  
  
thrust::sort_by_key(keys, keys + N, values);
```
- ```
// keys is now  { 1,  2,  4,  5,  7,  8}  
// values is now {'a', 'c', 'b', 'e', 'f', 'd'}
```

# Asynchronous operations

- Thrust also has some support for asynchronous copies and operations.
- Assign operations on streams: `thrust::device.on(stream)`
- Operations in the namespace `thrust::async`
- Capture operation in a `thrust::device_event`, and schedule operations by passing dependencies: `thrust::device.after(event)`.
- Capture results in `thrust::device_future`

# Asynchronous operations

- *// Asynchronously transfer to the device.*
- `thrust::device_vector<double> d_vec(h_vec.size());`
- `thrust::device_event e = thrust::async::copy(h_vec.begin(), h_vec.end(),  
d_vec.begin());`
- *// After the transfer completes, asynchronously compute the sum on the device.*
- `thrust::device_future<double> f0 = thrust::async::reduce(thrust::device.after(e),  
d_vec.begin(), d_vec.end(),  
0.0, thrust::plus<double>());`

# BLAS: Basic Linear Algebra Subroutines (BLAS)

- Basic functions that provide standard building blocks for matrix and vector operations.
- Classified into three classes:
  - BLAS1: Level 1 BLAS: Basic vector operations
  - BLAS2: Level 2 BLAS: Matrix-vector operations
  - BLAS3: Level 3 BLAS: Matrix-matrix operations
- Documentation and available operations: <https://www.netlib.org/blas/>

# BLAS: Basic Linear Algebra Subroutines (BLAS)

- Basic functions that provide standard building blocks for matrix and vector operations.
- Classified into three classes:

- BLAS1: Level 1 BLAS operations
  - Copy, scale, dot product, Norm etc
- Naming: <S: Single precision (32 bit)

D: Double precision (64 bit)

C: Single Complex (32 bit)

Z: Double Complex (64 bit)>

+ operation name

→ `scopy(x, y); // vector copy`

Generate plane rotation  
 Generate modified plane rotation  
 Apply plane rotation  
 Apply modified plane rotation  
 $x \leftrightarrow y$   
 $x \leftarrow \alpha x$   
 $y \leftarrow x$   
 $y \leftarrow \alpha x + y$   
 $dot \leftarrow x^T y$   
 $dot \leftarrow x^T y$   
 $dot \leftarrow x^H y$   
 $dot \leftarrow \alpha + x^T y$   
 $nrm2 \leftarrow \|x\|_2$   
 $asum \leftarrow \|re(x)\|_1 + \|im(x)\|_1$   
 $amax \leftarrow 1^{st} k \ni |re(x_k)| + |im(x_k)|$   
 $= \max(|re(x_i)| + |im(x_i)|)$

# BLAS: Basic Linear Algebra Subroutines (BLAS)

- BLAS2: Level 2 BLAS operations
  - Matrix vector product, rank operations
  - Naming: <S: Single precision (32 bit)

D: Double precision (64 bit)

C: Single Complex (32 bit)

Z: Double Complex (64 bit)>

+ matrix type (general, symmetric, banded  
etc) + operation name:

→ `sgemv(...); // matrix-vector`  
  
`// product`

$$\begin{aligned} y &\leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n \\ y &\leftarrow \alpha Ax + \beta y, y \leftarrow \alpha A^T x + \beta y, y \leftarrow \alpha A^H x + \beta y, A - m \times n \\ y &\leftarrow \alpha Ax + \beta y \\ y &\leftarrow \alpha Ax + \beta y \\ y &\leftarrow \alpha Ax + \beta y \\ y &\leftarrow \alpha Ax + \beta y \\ y &\leftarrow \alpha Ax + \beta y \\ y &\leftarrow \alpha Ax + \beta y \\ x &\leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x \\ x &\leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x \\ x &\leftarrow Ax, x \leftarrow A^T x, x \leftarrow A^H x \\ x &\leftarrow A^{-1}x, x \leftarrow A^{-T}x, x \leftarrow A^{-H}x \\ x &\leftarrow A^{-1}x, x \leftarrow A^{-T}x, x \leftarrow A^{-H}x \\ x &\leftarrow A^{-1}x, x \leftarrow A^{-T}x, x \leftarrow A^{-H}x \end{aligned}$$

$$\begin{aligned} A &\leftarrow \alpha xy^T + A, A - m \times n \\ A &\leftarrow \alpha xy^T + A, A - m \times n \\ A &\leftarrow \alpha xy^H + A, A - m \times n \\ A &\leftarrow \alpha xx^H + A \\ A &\leftarrow \alpha xx^H + A \\ A &\leftarrow \alpha xy^H + y(\alpha x)^H + A \\ A &\leftarrow \alpha xy^H + y(\alpha x)^H + A \\ A &\leftarrow \alpha xx^T + A \\ A &\leftarrow \alpha xx^T + A \\ A &\leftarrow \alpha xy^T + \alpha yx^T + A \\ A &\leftarrow \alpha xy^T + \alpha yx^T + A \end{aligned}$$

# BLAS: Basic Linear Algebra Subroutines (BLAS)

- BLAS3: Level 3 BLAS operations

- Matrix matrix operations, rank k operations

- <S: Single precision (32 bit)  $C \leftarrow \alpha op(A)op(B) + \beta C, op(X) = X, X^T, X^H, C - m \times n$

$$C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^T$$

D: Double precision (64 bit)

$$C \leftarrow \alpha AB + \beta C, C \leftarrow \alpha BA + \beta C, C - m \times n, A = A^H$$

$$C \leftarrow \alpha AA^T + \beta C, C \leftarrow \alpha A^T A + \beta C, C - n \times n$$

C: Single Complex (32 bit)

$$C \leftarrow \alpha AA^H + \beta C, C \leftarrow \alpha A^H A + \beta C, C - n \times n$$

$$C \leftarrow \alpha AB^T + \bar{\alpha} BA^T + \beta C, C \leftarrow \alpha A^T B + \bar{\alpha} B^T A + \beta C, C - n \times n$$

$$C \leftarrow \alpha AB^H + \bar{\alpha} BA^H + \beta C, C \leftarrow \alpha A^H B + \bar{\alpha} B^H A + \beta C, C - n \times n$$

Z: Double Complex (64 bit)>

$$B \leftarrow \alpha op(A)B, B \leftarrow \alpha Bop(A), op(A) = A, A^T, A^H, B - m \times n$$

$$B \leftarrow \alpha op(A^{-1})B, B \leftarrow \alpha Bop(A^{-1}), op(A) = A, A^T, A^H, B - m \times n$$

+ matrix type (general, symmetric, banded etc) +

operation name:

→ `sgemm(...); // matrix-matrix`

`// product`

# LAPACK

- Provides routines for solving linear systems, least-squares, eigenvalue problems, SVD and various factorizations (LU, Cholesky, QR etc)
- Leverages BLAS where possible.
- Reorganized algorithms to use block-matrix operations (gemm etc) for better memory accesses and overall higher throughput.
- Vendors (NVIDIA, Intel, AMD) implement highly tuned BLAS for their hardware.
- LAPACK can use the vendor-provided, standardized interface.

$$\begin{bmatrix} L & A & P & A & C & K \\ L & -A & P & -A & C & -K \\ L & A & P & A & -C & -K \\ L & -A & P & -A & -C & K \\ L & A & -P & -A & C & K \\ L & -A & -P & A & C & -K \end{bmatrix}$$



## cuBLAS $\rightarrow$ BLAS for NVIDIA GPUs


- Provides BLAS operations (Level 1, 2 and 3) for NVIDIA GPUs.
- Default storage is column-major and 1-based indexing.
  - Need to be careful when calling functions from C/C++.
- The general workflow:
  - Create a handle: A context to allow multi-threading and multi-GPU setups.
  - Setup data on the GPU in the required layout.
  - Call the cuBLAS API and assign it to a stream if necessary.
- Use the new API and not the legacy API: `#include <cublas_v2.h>`

# cuBLAS $\rightarrow$ BLAS for NVIDIA GPUs

- DEMO
- <https://github.com/NVIDIA/CUDALibrarySamples/tree/master/cuBLAS>

# Sparsity ?

- Consider a matrix of size  $(m \times n)$  most of whose elements are zeros.
- Storing matrices in dense requires  $\mathcal{O}(mn)$  elements.
  - Wasteful when most elements are zeros.
  - Sparsity ratio:  $\phi = \frac{nnz}{mn}$ , where  $nnz$  denotes the number of nonzeros in the matrix.
- We can do better and store only the nonzero elements.
  - Sparse formats: Specialized formats to store nonzeros and their locations.
  - Examples: COO, CSR, ELL, etc



# Sparse formats: COO

- COO: Coordinate format:
  - Store the row indices, column indices and values.

$$\begin{pmatrix} 1 & 0 & 0 & 2 \\ 0 & 3 & 0 & 0 \\ 0 & 4 & 5 & 0 \\ 6 & 0 & 0 & 7 \end{pmatrix}$$

# Sparse formats: COO

- COO: Coordinate format:
  - Store the row indices, column indices and values.

Row indices: [0 3 1 1 2 0 3]

Column indices: [0 0 1 2 2 3 3]

Values: [1 2 3 4 5 6 7]

|   | 0 | 1 | 2 | 3 |   |
|---|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 2 | 0 |
| 1 | 0 | 3 | 0 | 0 | 1 |
| 2 | 0 | 4 | 5 | 0 | 2 |
| 3 | 6 | 0 | 0 | 7 | 3 |

- Storage complexity:  $\mathcal{S} = 3nnz$

# Sparse formats: CSR

- CSR: Compressed sparse row
  - Store the column indices, row pointers and values.

$$\begin{array}{cccc} 0 & 1 & 2 & 3 \\ \left( \begin{array}{cccc} 1 & 0 & 0 & 2 \\ 0 & 3 & 0 & 0 \\ 0 & 4 & 5 & 0 \\ 6 & 0 & 0 & 7 \end{array} \right) & \begin{array}{c} 0 \\ 1 \\ 2 \\ 3 \end{array} \end{array}$$

# Sparse formats: CSR

- CSR: Compressed sparse row

- Store the column indices, row pointers and values.

Row pointers: [0 2 3 5 7]

Column indices: [0 0 1 2 2 3 3]

Values: [1 2 3 4 5 6 7]

Total number of  
nonzeros in  
matrix

Of size  $(n+1)$ ,  
with starting  
element 0

|   | 0 | 1 | 2 | 3 |   |
|---|---|---|---|---|---|
| 1 | 0 | 0 | 2 |   | 0 |
| 0 | 3 | 0 | 0 |   | 1 |
| 0 | 4 | 5 | 0 |   | 2 |
| 6 | 0 | 0 | 7 |   | 3 |

- Storage complexity:  $\mathcal{S} = 2nnz + n + 1$

# Sparse formats: ELL

- ELL: ELLPack
  - Store a fixed number of nonzeros in each row

$$\begin{array}{cccc|c}
 0 & 1 & 2 & 3 & \\
 \hline
 \left( \begin{array}{cccc}
 1 & 0 & 0 & 2 \\
 0 & 3 & 0 & 0 \\
 0 & 4 & 5 & 0 \\
 6 & 0 & 0 & 7
 \end{array} \right) & \begin{array}{c} 0 \\ 1 \\ 2 \\ 3 \end{array}
 \end{array}$$



# Sparse formats: ELL

- ELL: ELLPack
  - Store a fixed number of nonzeros in each row

Num nnz per row,  $l$ : 2

Column indices: [0 1 1 0 3 0 2 3]

Column-major  
storage

|   | 0 | 1 | 2 | 3 |
|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 2 |
| 1 | 0 | 3 | 0 | 0 |
| 2 | 0 | 4 | 5 | 0 |
| 3 | 6 | 0 | 0 | 7 |

Values: [1 3 4 6 2 0 5 7]

Stores an explicit  
zero

- Storage complexity:  $\mathcal{S} = 2(n \times l)$

# cuSPARSE → Sparse BLAS for NVIDIA GPUs

- Provides BLAS operations for sparse matrices on NVIDIA GPUs.
- Choosing the correct sparse format and ensuring data in that layout is user responsibility.
- The general workflow:
  - Create a handle: A context to allow multi-threading and multi-GPU setups.
  - Setup data on the GPU in the required format and layout.
  - Call the cuSPARSE API and assign it to a stream if necessary.
- `#include <cusparse.h>`

# cuSPARSE $\rightarrow$ Sparse BLAS for NVIDIA GPUs

- DEMO
- <https://github.com/NVIDIA/CUDALibrarySamples/tree/master/cuSPARSE>

# Tensors

- Higher-order arrays. Generalizes the matrix concept in higher dimensions.
  - Scalar  $\rightarrow$  order 0 tensor
  - Vector  $\rightarrow$  order 1 tensor
  - Matrix  $\rightarrow$  order 2 tensor etc.
- An order- $n$  tensor has  $n$  modes. Each mode has an extent (size in that dimension) and a stride.
- Einstein notation:  $y = \sum_i x^i z_i$  is represented as  $x^i z_i$ . With repeated indices, the summation is assumed to be implicit.
- Example of a tensor operation:  $C_{a,b,c} = A_{a,k,c} B_{k,b}$

# cuTENSOR → Tensor operations on NVIDIA GPUs

- Provides tensor operations for NVIDIA GPUs: Tensor contraction, reduction, and element-wise operations
- Mixed-precision support to utilize tensor cores.
- The general workflow:
  - Create a handle: A context to allow multi-threading and multi-GPU setups.
  - Setup data on the GPU in the required format and layout: More involved than matrix/vectors
  - Setup “Plan” cache to enable efficient memory usage to avoid re-allocations.
  - Call the cuTENSOR API and assign it to a stream if necessary.
- `#include <cutensor.h>`

# cuTENSOR $\rightarrow$ Tensor operations on NVIDIA GPUs

- DEMO
- <https://github.com/NVIDIA/CUDALibrarySamples/tree/master/cuTENSOR>

# Computational physics

Equation describing physics



Non-linear discretization



Linearization of the  
non-linear iteration



Each step requires a  
solution of a coupled  
linear system of the form

$$\mathbf{A}\mathbf{X} = \mathbf{B}$$

$$\phi'(t) = \underbrace{R(t, \phi(t))}_{\text{Reaction term}} + \underbrace{F(t, \phi(t))}_{\text{Forcing term}}$$



$$\begin{bmatrix} A_{11} & \cdots & \cdots & A_{1K} \\ \vdots & A_{22} & & \vdots \\ \vdots & & \ddots & \\ A_{K1} & \cdots & & A_{KK} \end{bmatrix} \begin{bmatrix} \Phi_1 \\ \Phi_2 \\ \vdots \\ \Phi_K \end{bmatrix} = \begin{bmatrix} G_1 \\ G_2 \\ \vdots \\ G_K \end{bmatrix}$$

# Machine learning and optimization

- Principal Component Analysis (PCA): Identify the main components in high-dimensional data → Uses an eigenvalue solver.
  - Singular value decomposition (SVD):  $M = U\Sigma V^*$
  - Also used for data compression, noise identification etc.
- Regression/fitting/supervised machine learning:  $\min_{x \in \mathbb{R}^n} \|Ax - b\|^2$



## cuSolver: LAPACK on GPUs

- Using cuBLAS and cuSPARSE, cuSOLVER provides LAPACK type routines for dense and sparse data structures on GPUs.
- cuSolverDN: Dense LAPACK: Factorization, eigenvalue solvers etc.
- cuSolverSP: Sparse LAPACK: Factorizations and eigenvalue solvers for sparse matrices stored in CSR format.
- DEMO (<https://github.com/NVIDIA/CUDALibrarySamples/tree/master/cuSOLVER>)

# Summary

- cuBLAS: BLAS operations for dense matrices
- cuSPARSE: BLAS operations for Sparse matrices
- cuTENSOR: Tensor operations
- cuSOLVER: Solvers and LAPACK routines for NVIDIA GPUs

# Next lecture

- Distributed computing basics.
- Distributed programming models: MPI, OpenSHMEM, NVSHMEM
- Multi-GPU programming with CUDA.