

# Tutorial 2: Cluster access and Parallel Computing

## Informatik elective: GPU Computing

Pratik Nayak

Licensed under



# In this session

- Cluster credentials and access
  - Navigating the cluster and check specifications
- Cluster performance check
- Estimate cluster performance: Peak FLOP/s, memory bandwidth, etc
- Clone the main exercise repository.
  - Solve exercises

## [HANDS-ON] Login to cluster

- Login to the cluster:
  - Ensure you are logged into to eduVPN (you might face issues with eduroam)
  - `ssh -i <your_ssh_key> <username>@vm.ginkgo-project.de`
  - `ssh -i <your_ssh_key> <username>@10.152.225.226`
- `pwd` and check `/home/<username>` exists
- Does `nvidia-smi` show the available GPUs ?

## [HANDS-ON] Clone the exercise framework repo

- We will use [gitlab.lrz.de](https://gitlab.lrz.de) as the remote repository for the exercises.
- Create an account on [gitlab.lrz.de](https://gitlab.lrz.de)
  - Login and update username
  - Tell me your username and I can add you to the group.
- Clone the repository:

```
git clone https://gitlab.lrz.de/2024ws-gpu-computing/exercises.git
```

## [HANDS-ON] Run the hello world GPU example

- Build the code with CMake
- CUDA and CMake should already be available (`cmake --version` and `nvcc --version`).
- Regular build with CMake should link with CUDA and create an executable (See previous tutorial slides).
- Run the executable and check if you get output information on the current GPU.

## [HANDS-ON] Arithmetic intensity

- Look at the example for the dot product, and implement the other operations:
  - axpy:  $y = y + \alpha x$ , where  $y$  and  $x$  are arrays and  $\alpha$  is a scalar.
  - matvec:  $x = A*b$ , where  $A$  is a matrix and  $b$  is a vector.
  - matmul:  $C = A*B$ , where  $A$  and  $B$  are matrices
- Measure the performance in GFlops/s of these operations (with increasing input sizes) and plot a roofline on your local system, and also on the cluster.
- Are these operations memory bound/compute bound ?